

Beyond Classical Search

Unit-II

Local Search Algorithms

- The informed and uninformed search expands the nodes systematically in two ways:
 - keeping different paths in the memory and
 - selecting the best suitable path,
- In “local search algorithms”, the path cost does not matter, and only focus on solution-state needed to reach to the **goal node**.
- A local search algorithm completes its task by traversing on a single current node rather than multiple paths and following the neighbors of that node.
- In many optimization problems, the path to the goal is **irrelevant**; the goal state itself is the solution.

Local Search Algorithms

- ***Advantages:***
 1. Local search algorithms use a very little or constant amount of memory.
 2. They find a reasonable solution in large or infinite state spaces where the classical or systematic algorithms do not work suitably.
- Local search algorithms can also solve **optimization problems**, in which the aim is to find the best state according to an objective function.

Local Search Algorithms...

- The local search algorithm is also work for a pure optimized problem, where a **pure optimization problem** is one where all the nodes can give a solution.
- But the target is to find the best state out of all according to the objective function.
- An objective function is a function whose value is either minimized or maximized in different contexts of the optimization problems.

Local Search Algorithms...

- To understand the local search algorithms, consider the below state-space landscape diagram having:
- **Location:** It is defined by the state.
- **Elevation:** It is defined by the value of the objective function or heuristic cost function.


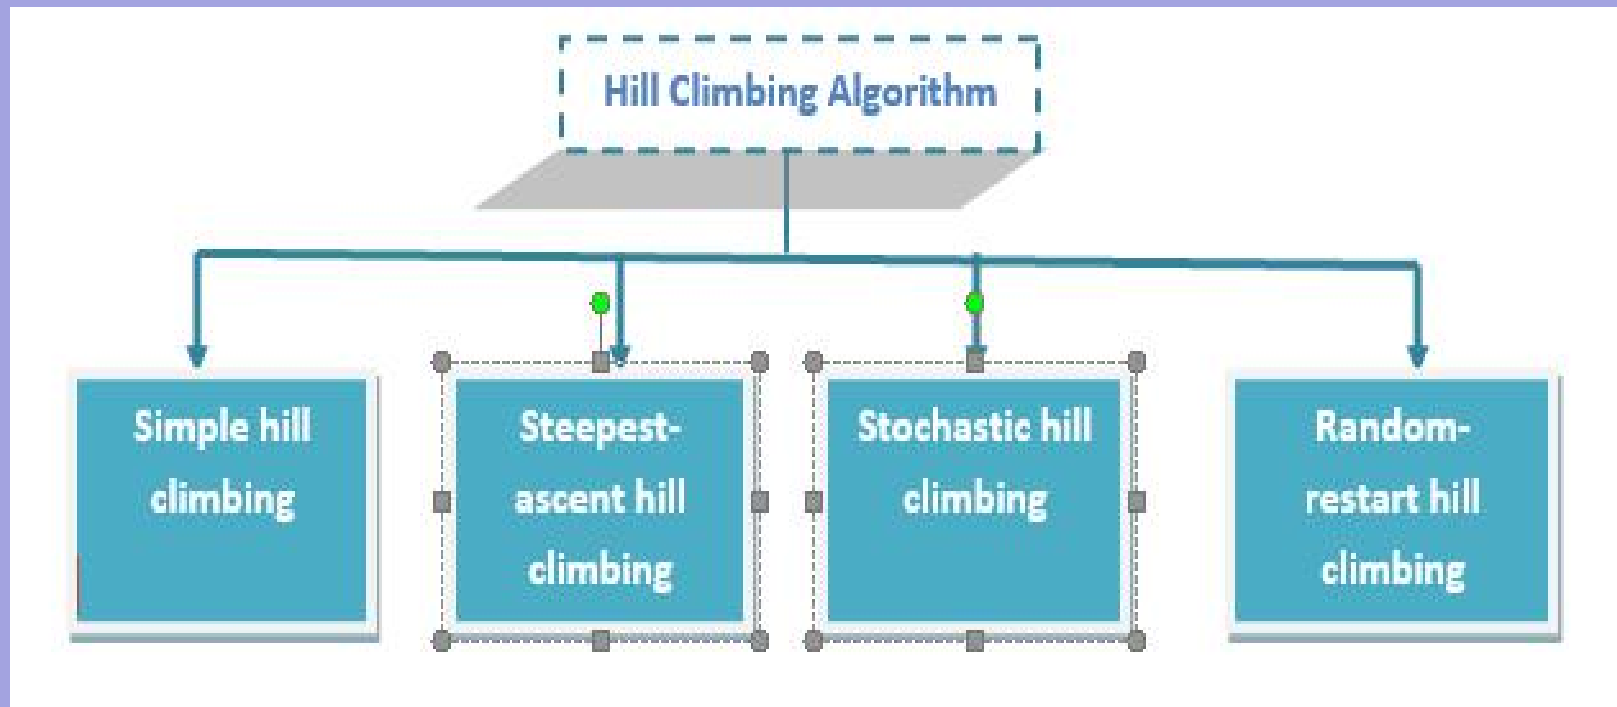


Figure: A one-dimensional state-space landscape in which elevation corresponds to the objective function. The aim is to find the global maximum.

Local Search Algorithms...

- The local search algorithm explores the above landscape by finding the following two points:
- **Global Minima:** If the elevation corresponds to the cost, then the task is to find the lowest valley, which is known as Global Minimum.
- **Global Maxima:** If the elevation corresponds to an objective function, then it finds the highest peak which is called as Global Maxima. It is the highest point in the valley.
- Different types of local searches:
 - Hill-climbing Search
 - Simulated Annealing
 - Local Beam Search

Local Search Algorithms...



Hill climbing

- It is often used when a **good heuristic function** is available for evaluating states.
- But when no other useful knowledge is available. This algorithm is simply a loop that continuously moves in the direction of increasing value i.e., uphill.
- It terminates when it reaches a “peak” where no neighbor has a higher value.
- **The algorithm doesn't maintain a search tree, so the current node data structure only records the state and its objective function value.**
- Hill – climbing doesn't look ahead beyond the immediate neighbors of the current state.

Simple Hill climbing algorithm

1. Evaluate the initial state (IS). If it is the goal state (GS) , then return it and quit. Else consider IS as the current state (CS) and proceed.
2. Loop until a solution is found or there are no new operator (OP) to be applied to the CS.
 - a) Select an OP that has not yet been applied to the CS and apply it to produce a new state (NS).
 - b) Evaluate the NS:
 - If NS is a GS , then return it and quit.
 - If it is not a GS but better than the CS, then consider it as the current state (CS) and proceed.
 - If NS is not better than CS then continue in the loop by selecting the next appropriate OP for CS.

Steepest – Ascent Hill climbing algorithm

- It considers all the moves from the CS and selects the best one as the next state.
- It is also known as a **greedy approach**
- It is also called **gradient search (or gradient ascent/descent)**.

Algorithm:

1. Evaluate the initial state (IS). If it is the goal state (GS) , then return it and quit. Else consider IS as the current state (CS) and proceed.
2. Loop until a solution is found or until a complete iteration produces no change to the CS:

Steepest – Ascent Hill climbing algorithm

- a) Let successor (SUC) be a state such that any NS that can be generated from CS is better than SUC. (i.e., setting SUC to a minimum value at the beginning of an iteration or set CS as SUC)
- b) For each operator OP that applies to the CS do:
 - I. Apply OP to CS and generate a NS.
 - II. Evaluate the NS. If it is a GS then return it and quit. If not , compare it with SUC. If NS is better than SUC, then set SUC to NS; else leave SUC unchanged.
- c) If the SUC is better than CS, then set CS to SUC (i.e., move to the next best state)

Steepest Hill-climbing search

- "Like climbing Everest in thick fog with amnesia"

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node


  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
```

Stochastic hill climbing search

```
current position = initial solution;  
repeat  
  for All neighbours of current position do  
    Obtain a random neighbour;  
    if cost of neighbour  $\leq$  cost of current position then  
      current position = neighbour position;  
      break;  
    end  
  end  
until cost of current position  $\leq$  cost of all its neighbours;
```

Random-restart hill climbing search

- Random-restart algorithm is based on try and try strategy.
- It iteratively searches the node and selects the best one at each step until the goal is not found.
- The success depends most commonly on the shape of the hill.
- If there are few plateaus, local maxima, and ridges, it becomes easy to reach the destination.




Limitations: Hill-climbing search

- Most of the hill climbing search algorithms may fail to find a optimum solution.
- Either algorithm may terminate not by finding a goal state but by getting to a state from which no better states can be generated.
- Hill Climbing is not **complete**; Unless we introduce backtracking
- Hill Climbing is not **optimal**; Solution found is a local optimum
- This will happen if the program has reached either a **local maximum, a plateau or a ridge**.


Limitations: Hill-climbing search

- A **local maximum** is a state that is better than all its neighbors but is not better than some other states further away.
- Solution of local maxima:-
 - Move in some arbitrary direction
 - Back track to an ancestor and try some other alternatives.




Limitations: Hill-climbing search

- **A plateau** is a flat area in the search space in which all the neighboring states have the same heuristic function value.
- Solution of plateau
 - Expand few generation ahead to move to a different section of the search space.




Limitations: Hill-climbing search

- **A ridge** is an area in the search space which is higher than its surroundings but itself has slopes.
- It is not possible to traverse a ridge by a single move i.e., no such operator is available.
- **Solution of ridges:-**
 - Apply several operators before doing the evaluation




Limitations: Hill-climbing search

- Problem: depending on initial state, can get stuck in local maxima



Hill-climbing search: n-queens problem

- **Objective/Goal:** Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal which means no queen attacking another.



- **States:** n queens on board, one per column
- **Actions:** move a queen in its column
- **Heuristic value function:** number of conflicts
- Move a queen to reduce number of conflicts


Hill-climbing search: 8-queens problem

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♙	13	16	13	16
♙	14	17	15	♙	14	16	16
17	♙	16	18	15	♙	15	♙
18	14	♙	15	15	14	♙	16
14	14	13	17	12	14	12	18


- h = number of pairs of queens that are attacking each other, either directly or indirectly.
- $h = 17$ for the above state.
- Therefore the best greedy move is to move a queen to a square labeled with 12.

Hill-climbing search: 8-queens problem

- What is h value here? Is it global minimum?

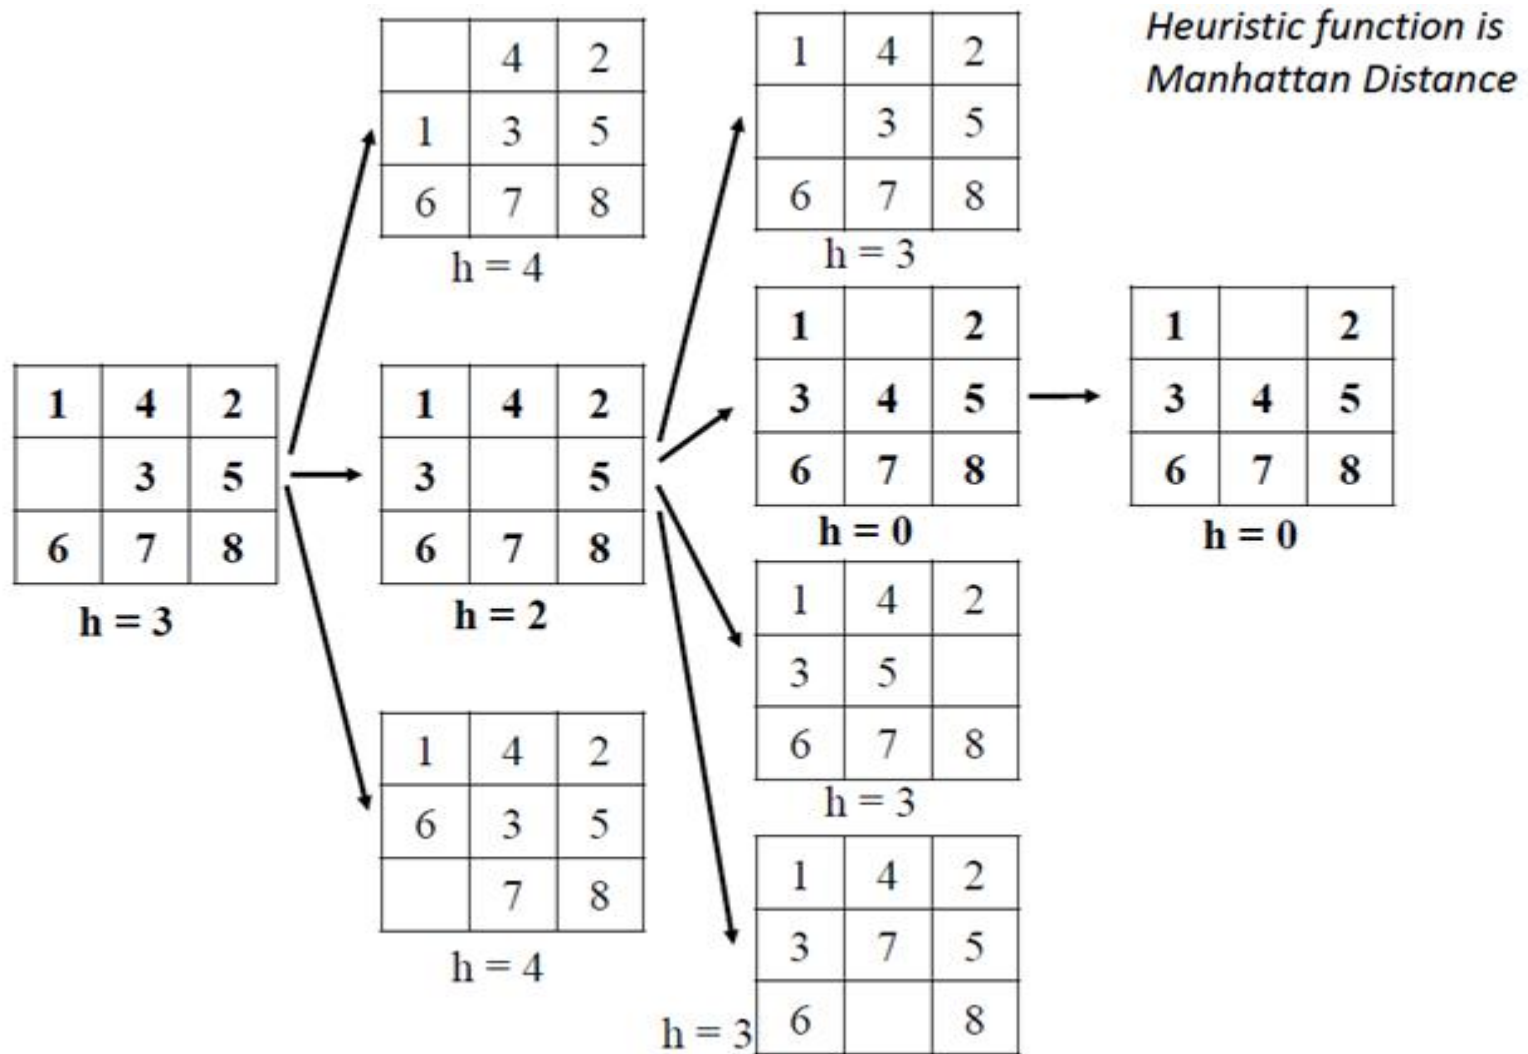


Hill-climbing search: 8-queens problem

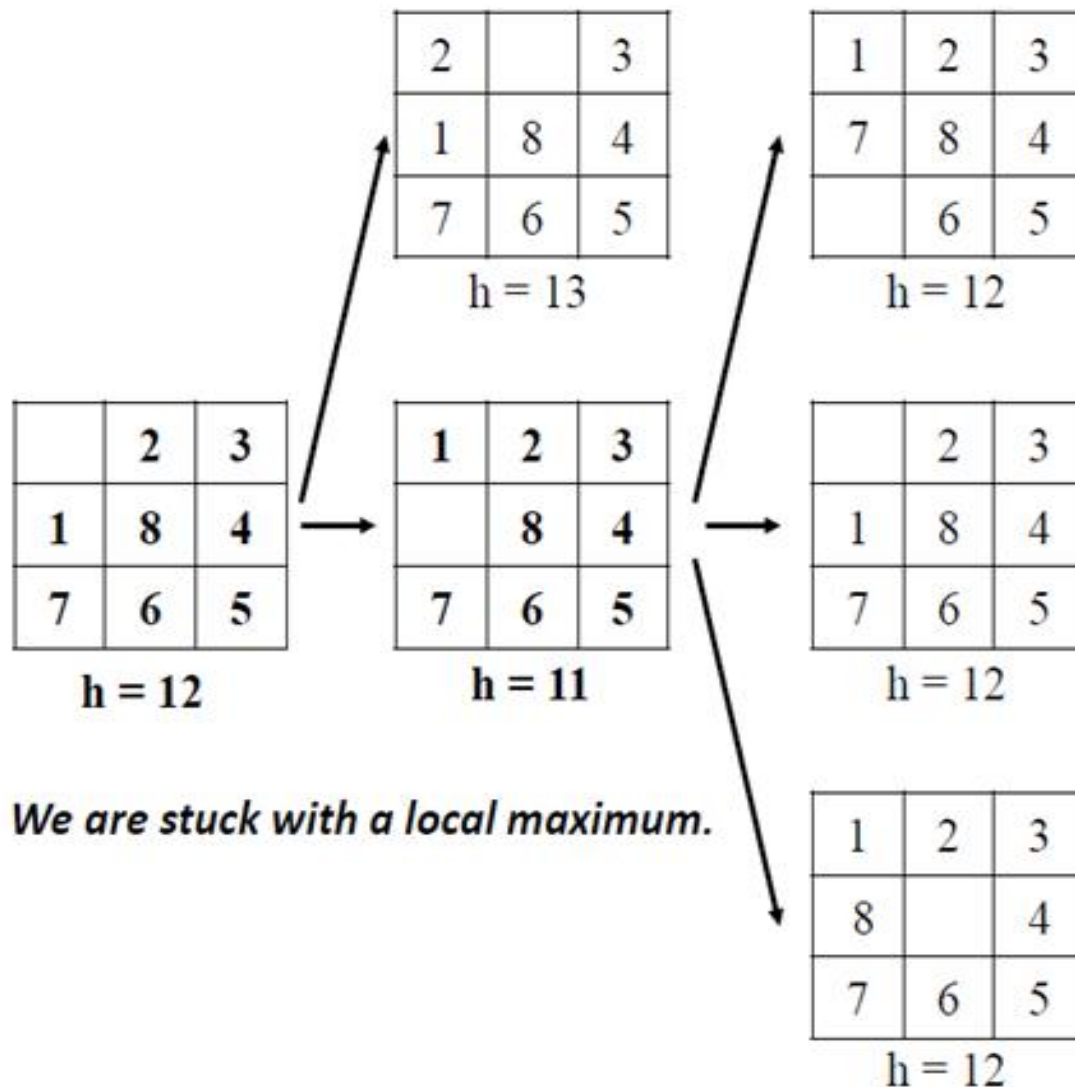


- A local minimum with $h = 1$

Hill-climbing search: 8-puzzle problem



Hill-climbing example: 8-puzzle problem

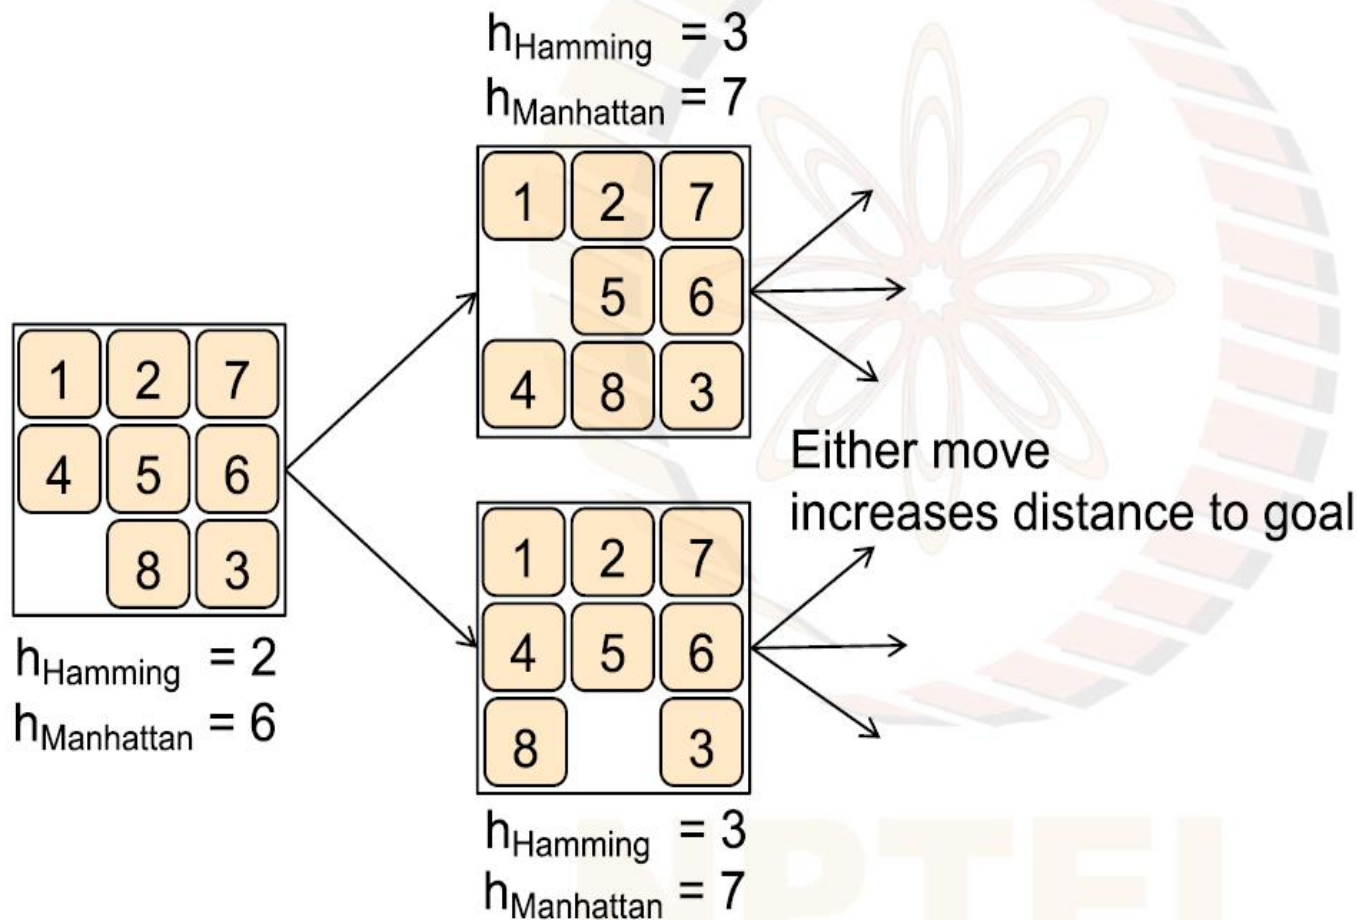


*Heuristic function is
Manhattan Distance*

We are stuck with a local maximum.


Hill-climbing example: 8-puzzle problem

8-puzzle: A local minimum



Hill-climbing example: Traveling Salesman Problem(TSP)

- Find the shortest tour connecting n cities
- **State space:** all possible tours
- **Objective function:** length of tour
- What's a possible local improvement strategy?
 - Start with any complete tour, perform pairwise exchanges



Simulated Annealing

- A **hill-climbing algorithm** that never makes “downhill” moves toward states with lower value (or higher cost) is guaranteed to be **incomplete**, because *it can get stuck on a local maximum*.
 - In contrast, a **purely random walk**—that is, moving to a successor chosen uniformly at random from the set of successors—is **complete but extremely inefficient**.
 - Therefore, it seems reasonable to *combine hill climbing with a random walk in some way that yields both efficiency and completeness*.
- Simulated Annealing (SA) is applied to solve optimization problems.
- SA is a **stochastic algorithm**.
- Simulated Annealing is allow moves to **inferior neighbors** with a probability that is regulated over time.
- Escape local maxima by allowing **some "bad" moves** but gradually decrease their probability.

Simulated Annealing

- The probability is controlled by a parameter called **temperature**
- Higher temperatures allow more bad moves than lower temperatures
- Annealing: Lowering the temperature gradually
- Quenching: Lowering the temperature rapidly .
- **It can be proven that:** If T decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1.
- **SA** is motivated by the physical annealing process.
- Method proposed in 1983 by IBM researchers for solving VLSI layout problems (Kirkpatrick et al, *Science*, 220:671-680, 1983). It will always find the global optimum
- **Other applications:** Traveling salesman, Graph partitioning, Graph coloring, Scheduling, Facility Layout, Image Processing, useful for some problems, but can be very slow.

Simulated Annealing

function SIMULATED-ANNEALING(*problem*, *schedule*) **return** a solution state

input: *problem*, a problem

schedule, a mapping from time to temperature

local variables: *current*, a node.

next, a node.

T, a “temperature” controlling the prob. of downward steps

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

for *t* \leftarrow 1 to ∞ **do**

T \leftarrow *schedule*[*t*]

if *T* = 0 **then return** *current*


next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ VALUE[*next*] - VALUE[*current*]

if $\Delta E < 0$ **then** *current* \leftarrow *next*


else *current* \leftarrow *next* only with probability $e^{-\Delta E/T}$

Simulated Annealing




Simulated Annealing

$$\text{Probability of making a bad move} = e^{-\Delta E/T} = \frac{1}{e^{\Delta E/T}}$$



Since $\Delta E_1 > \Delta E_2$ moving from A to C is exponentially more probable than moving from A to B

Simulated Annealing

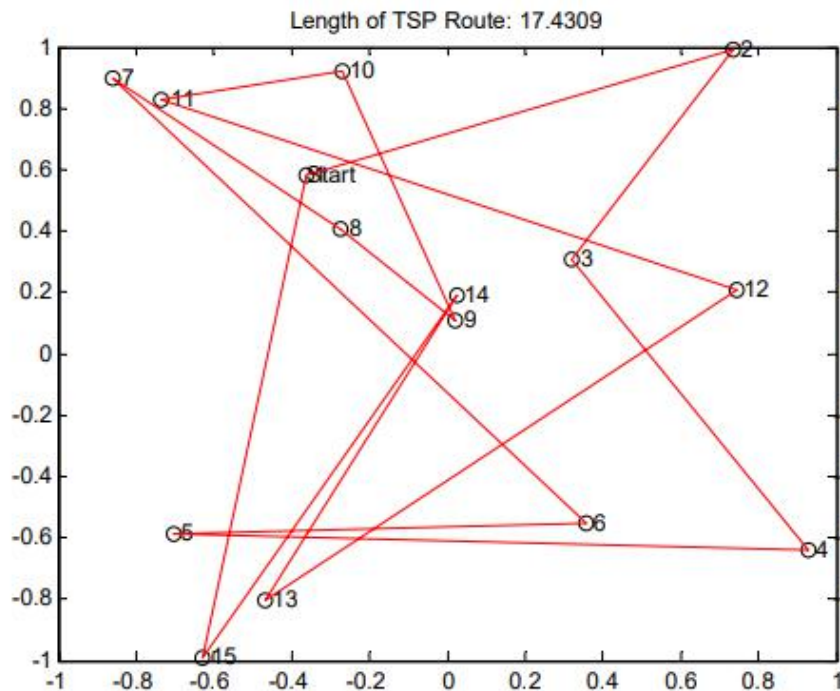


Simulated Annealing

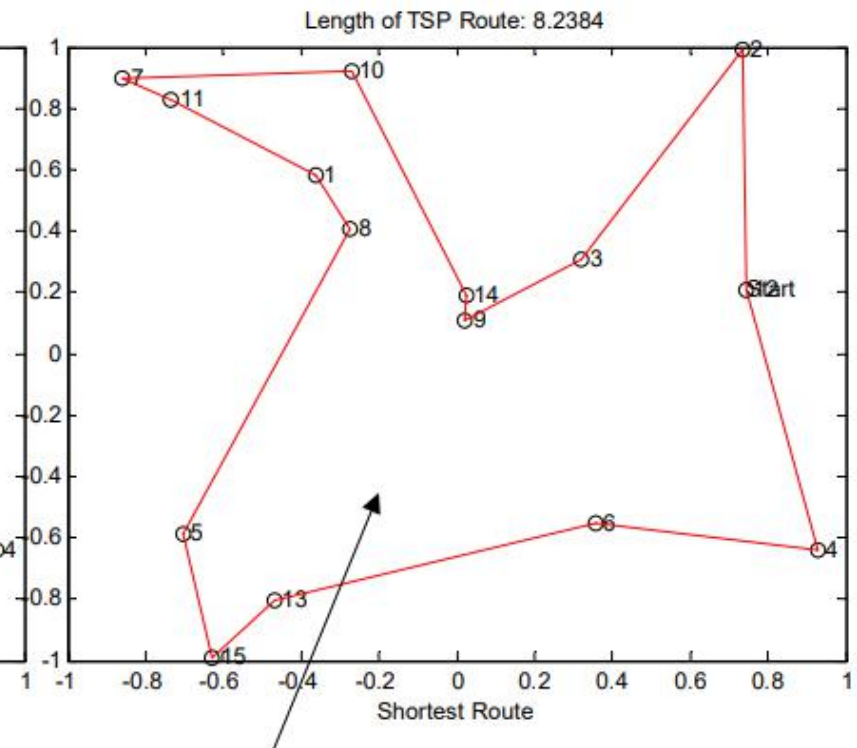
- The SA algorithm is based on the annealing process used in metallurgy, where a metal is heated to a high temperature quickly and then gradually cooled.
- At high temperatures, the atoms move fast, and when the temperature is reduced, their kinetic energy decreases as well.
- At the end of the annealing process, the atoms fall into a more ordered state, and the material is more ductile and easier to work with.
- Similarly, in SA, a search process starts with a high-energy state (an initial solution) and gradually lowers the temperature (a control parameter) until it reaches a state of minimum energy (the optimal solution).

Simulated Annealing

Initial (Random) Route
Length: 17.43



Final (Optimized) Route
Length: 8.24



Result with SA

Local beam Search

- Keeping just one node in memory might seem an extreme reaction to the problem of memory limitation, but this is what we do in local searches and in simulated annealing.
- Local Beam Search is a **heuristic search algorithm** used for optimization problems.
- Beam search is an optimization of **best-first search** that reduces its memory requirements.
- It is used to explore a large search space efficiently by focusing on a subset of the search space.
- The main advantage of local beam search is that it can explore a large search space efficiently by focusing on the best subset of solutions.
- This can help to avoid getting stuck in local optima and improve the quality of the solutions found


Local beam Search: Algorithm

```
function BEAM-SEARCH(problem, k) returns a solution state
  start with k randomly generated states
  loop
    generate all successors of all k states
    if any of them is a solution then return it
    else select the k best successors
```

Local beam Search...

- **Idea:** Keeping only one node in memory is an extreme reaction to memory problems.
- Keep track of k states instead of one
- Initially: k randomly selected states
- Next: determine all successors of k states
- If any of successors is goal->finished
- Else select k best from successors and repeat.
- **E x a m p l e :**
<https://www.codecademy.com/resources/docs/ai/search-algorithms/beam-search>


Local beam Search



Continue till the goal state is found

Local beam Search:Example


Example-1
Beam= 2



Frontier List	Explored List
AB	S
CE	SAB
HJ	SABCE
LN	SABCEHJ
	SABCEHJLN

Local beam Search...

- Question:** Perform Local Beam Search With 2 Beams On The Graph With Edge Costs And Heuristic Values Below. Initial Nodes: A, H Goal Node: J



Stochastic Local beam Search


- Local beam search is not the same as *k random-start searches run in parallel!*
- Searches that find good states recruit other searches to join them
- **Problem:** quite often, all *k states end up on same local hill*
- **Idea: Stochastic beam search**
- Choose *k successors randomly, biased towards good ones.*
- Observe the close analogy to natural selection!
- Instead of choosing the best *k* from the pool of candidate successors, **stochastic beam search** chooses *k* successors at random, with the probability of choosing a given successor being an increasing function of its value.
- **Stochastic beam search** bears some resemblance to the process of natural selection, whereby the “successors” (offspring) of a “state” (organism) populate the next generation according to its “value” (fitness).

Genetic algorithm

- A **genetic algorithm (GA)** is a variant of stochastic beam search in which successor states are generated by combining two parent states rather than by modifying a single state.
- Start with k randomly generated states (**population**)
- A state or individual, is represented as a string over a finite alphabet (often a string of 0s and 1s)
- Each state is rated by an objective function, or (in GA terminology) the **fitness function**/evaluation function (fitness function). Higher values=Better results
- Pairs of individuals are randomly selected for **reproduction**.
- Produce the next generation of states by **Selection, Crossover, and Mutation**.


Genetic algorithm: Example

- An 8-queens state must specify the positions of 8 queens, each in a column of 8 squares, and so it requires $8 \times \log_2 8 = 24$ bits.
- Alternatively, the state could be represented as 8 digits, each in the range from 1 to 8.
- A state can be represented using a 8 digit string.
- Each digit in the range from 1 to 8 to indicate the position of the queen in that column.



Genetic algorithm: Example


- Crossover helps if substrings are meaningful components.
- The shaded columns are lost in the crossover step and the unshaded columns are retained





fitness:
#non-attacking queens

probability of being
regenerated
in next generation



- Fitness function: number of non-attacking pairs of queens (min = 0, max = $8 \times 7/2 = 28$)
- $24/(24+23+20+11) = 31\%$
- $23/(24+23+20+11) = 29\%$ etc

Genetic algorithm: Example

- In above example, the **initial population** has 4 states.
- A **fitness function** should return higher values for better states, so, for the 8-queens problem we use the number of non-attacking pairs of queens, which has a value of 28 for a solution.
 - The values of the four states are 24, 23, 20, and 11.
- The **probability of being chosen for reproducing** is directly proportional to the fitness score.
- Two pairs are selected at random for reproduction, in accordance with the probabilities.
 - Notice that one individual is selected twice and one not at all.

Genetic algorithm: Example

- The **crossover points** are after third digit in first pair and after fifth digit in second pair.
- The first child of the first pair gets the first three digits from the first parent and the remaining digits from the second parent,
- whereas the second child gets the first three digits from the second parent and the rest from the first parent.
- One digit was **mutated** in the first, third, and fourth offspring.
- In the 8-queensproblem, this corresponds to choosing a queen at random and moving it to a random square in its column.

Genetic algorithm: Example

- **Positive points:**
 - Random exploration can find solutions that local search can't (via crossover primarily)
 - Appealing connection to human evolution
 - “neural” networks, and “genetic” algorithms are **metaphors!**
- **Negative points:**
 - Large number of “tunable” parameters
 - Difficult to replicate performance from one problem to another
 - Lack of good empirical studies comparing to simpler methods
 - Useful on some (small?) set of problems but no convincing evidence that GAs are better than hill-